

```

# %%capture
# #!unzip Datasets.zip

# from google.colab import drive

# # Mount the Google Drive
# drive.mount('/content/drive')

# %%capture
# !pip install datasets
# !pip install transformers
# !pip install librosa
# !pip install jiwer
# !pip install evaluate

import os
import datasets
import pandas as pd
from sklearn.model_selection import train_test_split
from datasets import Dataset

# Set paths
csv_path = "/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final.csv"
audio_folder = "/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-w"

# Load the CSV
df = pd.read_csv(csv_path)
df = pd.read_csv(csv_path)
# Ensure the column names match
df.columns = ["Filename", "Transcription"] # Rename columns if needed

# Append '.wav' to the file names
df['Filename'] = df['Filename'].apply(lambda x: f"{x}.wav")

# Add full paths to the audio files
df['file_path'] = df['Filename'].apply(lambda x: os.path.join(audio_folder, x))

# Verify that all audio files exist
missing_files = df[~df['file_path'].apply(os.path.exists)]
if not missing_files.empty:
    print("The following audio files are missing:")
    print(missing_files)
    raise FileNotFoundError("Some audio files listed in the CSV are missing in the folder.")

# Split into train (27) and test (3)
train_df, test_df = train_test_split(df, test_size=0.2, random_state=42)

# Save splits to CSV for reference
train_csv_path = "train_split.csv"
test_csv_path = "test_split.csv"
train_df.to_csv(train_csv_path, index=False)
test_df.to_csv(test_csv_path, index=False)

# Convert to HuggingFace Dataset format
train_dataset = Dataset.from_pandas(train_df)
test_dataset = Dataset.from_pandas(test_df)

# Save HuggingFace datasets
train_dataset_path = "train_dataset"
test_dataset_path = "test_dataset"
train_dataset.save_to_disk(train_dataset_path)
test_dataset.save_to_disk(test_dataset_path)

# Output
print(f"Train set saved to: {train_csv_path} and {train_dataset_path}")
print(f"Test set saved to: {test_csv_path} and {test_dataset_path}")

↗ /home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please
from .autonotebook import tqdm as notebook_tqdm
Saving the dataset (1/1 shards): 100%|██████████| 2600/2600 [00:00<00:00, 719149.99 examples/s]
Saving the dataset (1/1 shards): 100%|██████████| 650/650 [00:00<00:00, 371582.06 examples/s]Train set saved to: train_s
Test set saved to: test_split.csv and test_dataset

from datasets import load_from_disk

train_dataset = load_from_disk("train_dataset")
test_dataset = load_from_disk("test_dataset")

```

```
print(train_dataset)
print(test_dataset)
```

```
Dataset({
  features: ['Filename', 'Transcription', 'file_path', '__index_level_0__'],
  num_rows: 2600
})
Dataset({
  features: ['Filename', 'Transcription', 'file_path', '__index_level_0__'],
  num_rows: 650
})
```

```
from datasets import ClassLabel
import random
import pandas as pd
from IPython.display import display, HTML
```

```
def show_random_elements(dataset, num_examples=10):
    assert num_examples <= len(dataset), "Can't pick more elements than there are in the dataset."
    picks = []
    for _ in range(num_examples):
        pick = random.randint(0, len(dataset)-1)
        while pick in picks:
            pick = random.randint(0, len(dataset)-1)
        picks.append(pick)

    df = pd.DataFrame(dataset[picks])
    display(HTML(df.to_html()))
```

```
show_random_elements(train_dataset)
```

	Filename	Transcription	file_path	__index_level_0__
0	jhon-01.1_126.wav	عبدالرحیم با سبہ پرہیزگار تواریختن پس خط ارشاد یمن حضرت میر محمد ہمدانی	/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-waves/jhon-01.1_126.wav	1956
1	jhon-03_80.wav	تہہ باغ سلیمان ناو پٹھ چھہ امیر لہجہ ہڑ مغازی البنی رسالہ ظل و بلبیل تہ	/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-waves/jhon-03_80.wav	2901
2	ishrat1-15_10.wav	امیر کنہی آمت توے پٹھ نانہ ٹھٹھاں از مہر بکسم وجاں شامن	/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-waves/ishrat1-15_10.wav	1480
3	jhon-01.2_30.wav	سیدیچ نہ تہ کار گر آتھی دوران و پٹھ زبیا سٹکھ ماجہ خاں بس ملکر	/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-waves/jhon-01.2_30.wav	2533
4	8140192.wav	گوہان یہ چھہ آسی ماتھہ ز عروج نیاد	/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-waves/8140192.wav	191
5	rafiya-04_16.wav	تہ شوی آتہ تہ دہرائڑ روکی پاٹھی قلمہ مہر شرانہ وسان	/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-waves/rafiya-04_16.wav	3090
6	jhon-01.1_87.wav	ایہ و اتان توم زن چھہ ساری کشیر اکھ زبٹھ ز پوٹھ سرٹنکر	/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-waves/jhon-01.1_87.wav	2496
7	ishrat1-34_5.wav	مشتھہ گوے کیا ز و تم گام آخر	/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-waves/ishrat1-34_5.wav	1710
8	jhon-03_43.wav	بن وجہہ شمس الضحیٰ من خدہ بدر الدجی	/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-waves/jhon-03_43.wav	2860
9	rafiya-05_23.wav	دلہیہ کردار حالہت بلکہ ملک تہ مادان جتہ پرتھہ کیلہہ پٹھ	/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-waves/rafiya-05_23.wav	3168

If there are any unwanted special characters in the dataset, we can remove them here, since there are none, I am keeping that as it is.

```
def extract_all_chars(batch):
    all_text = " ".join(batch["Transcription"])
    vocab = list(set(all_text))
    return {"vocab": [vocab], "all_text": [all_text]}
```

```
vocab_train = train_dataset.map(extract_all_chars, batched=True, batch_size=-1, keep_in_memory=True, remove_columns=train_dataset.get_column_names())
vocab_test = test_dataset.map(extract_all_chars, batched=True, batch_size=-1, keep_in_memory=True, remove_columns=test_dataset.get_column_names())
```

```
Map: 100%|██████████| 2600/2600 [00:00<00:00, 214961.08 examples/s]
Map: 100%|██████████| 650/650 [00:00<00:00, 251805.45 examples/s]
```

```
vocab_list = list(set(vocab_train["vocab"][0]) | set(vocab_test["vocab"][0]))
```

```
vocab_dict = {v: k for k, v in enumerate(sorted(vocab_list))}
vocab_dict
```

Show hidden output

```

vocab_dict["|"] = vocab_dict[" "]
del vocab_dict[" "]

vocab_dict["[UNK]"] = len(vocab_dict)
vocab_dict["[PAD]"] = len(vocab_dict)
len(vocab_dict)

↩ 60

import json
with open('vocab.json', 'w') as vocab_file:
    json.dump(vocab_dict, vocab_file)

from transformers import Wav2Vec2CTCTokenizer

tokenizer = Wav2Vec2CTCTokenizer.from_pretrained("./", unk_token="[UNK]", pad_token="[PAD]", word_delimiter_token="|", clear

from transformers import Wav2Vec2FeatureExtractor

feature_extractor = Wav2Vec2FeatureExtractor(feature_size=1, sampling_rate=16000, padding_value=0.0, do_normalize=True, retu

from transformers import Wav2Vec2Processor

processor = Wav2Vec2Processor(feature_extractor=feature_extractor, tokenizer=tokenizer)

train_dataset[0]["file_path"]

↩ ' /home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/final-
waves/farhat-03_21.wav'

```

Replacing the File Path with Actual Audio.

```

from datasets import load_from_disk, Audio

# Load datasets
train_dataset = load_from_disk("train_dataset") # Adjust to your actual path
test_dataset = load_from_disk("test_dataset")

# Rename 'file_path' to 'audio'
train_dataset = train_dataset.rename_column("file_path", "audio")
test_dataset = test_dataset.rename_column("file_path", "audio")

# # Cast the 'audio' column to use the Audio feature
train_dataset = train_dataset.cast_column("audio", Audio(sampling_rate=16_000))
test_dataset = test_dataset.cast_column("audio", Audio(sampling_rate=16_000))

# # Drop unnecessary columns if needed
train_dataset = train_dataset.remove_columns(["__index_level_0__"])
test_dataset = test_dataset.remove_columns(["__index_level_0__"])

# # Verify the dataset structure
print(train_dataset)
print(test_dataset)

# # Inspect the first example
print(train_dataset[0])

↩ Dataset({
  features: ['Filename', 'Transcription', 'audio'],
  num_rows: 2600
})
Dataset({
  features: ['Filename', 'Transcription', 'audio'],
  num_rows: 650
})
{'Filename': 'farhat-03_21.wav', 'Transcription': 'زأنم شاه صأبن بدشابس\مرا قېم پتم تم ييئلم', 'audio': {'path': '/ho
0.0256958 , 0.02392578}], 'sampling_rate': 16000}}

#print(test_dataset[0]['audio'])

rand_int = random.randint(0, len(train_dataset))

print("Target text:", train_dataset[rand_int]["Transcription"])
print("Input array shape:", train_dataset[rand_int]["audio"]["array"].shape)
print("Sampling rate:", train_dataset[rand_int]["audio"]["sampling_rate"])

```

```

➦ Target text: ژور مآری مئندی کأشیری چها شعر کأشیر زبأنی توش تم بوش عطا کران
Input array shape: (139205,)
Sampling rate: 16000

def prepare_dataset(batch):
    audio = batch["audio"]

    # batched output is "un-batched"
    batch["input_values"] = processor(audio["array"], sampling_rate=audio["sampling_rate"]).input_values[0]
    batch["input_length"] = len(batch["input_values"])

    batch["labels"] = processor(text=batch["Transcription"]).input_ids

    return batch

train_dataset = train_dataset.map(prepare_dataset, remove_columns=train_dataset.column_names)
test_dataset = test_dataset.map(prepare_dataset, remove_columns=test_dataset.column_names)

import torch

from dataclasses import dataclass, field
from typing import Any, Dict, List, Optional, Union

@dataclass
class DataCollatorCTCWithPadding:
    """
    Data collator that will dynamically pad the inputs received.
    Args:
        processor (:class:`~transformers.Wav2Vec2Processor`)
            The processor used for processing the data.
        padding (:obj:`bool`, :obj:`str` or :class:`~transformers.tokenization_utils_base.PaddingStrategy`, `optional`, default: `bool`)
            Select a strategy to pad the returned sequences (according to the model's padding side and padding index)
            among:
            * :obj:`True` or :obj:`'longest'`: Pad to the longest sequence in the batch (or no padding if only a single
              sequence if provided).
            * :obj:`'max_length'`: Pad to a maximum length specified with the argument :obj:`max_length` or to the
              maximum acceptable input length for the model if that argument is not provided.
            * :obj:`False` or :obj:`'do_not_pad'` (default): No padding (i.e., can output a batch with sequences of
              different lengths).
    """

    processor: Wav2Vec2Processor
    padding: Union[bool, str] = True

    def __call__(self, features: List[Dict[str, Union[List[int], torch.Tensor]]]) -> Dict[str, torch.Tensor]:
        # split inputs and labels since they have to be of different lengths and need
        # different padding methods
        input_features = [{"input_values": feature["input_values"]} for feature in features]
        label_features = [{"input_ids": feature["labels"]} for feature in features]

        batch = self.processor.pad(
            input_features,
            padding=self.padding,
            return_tensors="pt",
        )

        with self.processor.as_target_processor():
            labels_batch = self.processor.pad(
                label_features,
                padding=self.padding,
                return_tensors="pt",
            )

        # replace padding with -100 to ignore loss correctly
        labels = labels_batch["input_ids"].masked_fill(labels_batch.attention_mask.ne(1), -100)

        batch["labels"] = labels

        return batch

data_collator = DataCollatorCTCWithPadding(processor=processor, padding=True)

import evaluate

wer_metric = evaluate.load("wer")

➦ 2025-04-06 07:08:16.786995: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register cuFFT factory
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1743903496.837447 2721 cuda_dnn.cc:8310] Unable to register cuDNN factory: Attempting to register factory
E0000 00:00:1743903496.852595 2721 cuda_blas.cc:1418] Unable to register cuBLAS factory: Attempting to register facto

```

2025-04-06 07:08:16.968900: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to
 To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler fl
 Using the latest cached version of the module from /home/muzaffar/.cache/huggingface/modules/evaluate_modules/metrics/ev

```
from evaluate import load
```

```
cer_metric = load("cer")
```

➦ Using the latest cached version of the module from /home/muzaffar/.cache/huggingface/modules/evaluate_modules/metrics/ev

includes both WER and CER

```
def compute_metrics(pred):
    pred_logits = pred.predictions
    pred_ids = np.argmax(pred_logits, axis=-1)

    # Replace padding token (-100) with pad_token_id
    pred.label_ids[pred.label_ids == -100] = processor.tokenizer.pad_token_id

    # Decode predictions and labels to strings
    pred_str = processor.batch_decode(pred_ids)
    label_str = processor.batch_decode(pred.label_ids, group_tokens=False)

    if isinstance(label_str, list):
        if isinstance(pred_str, list) and len(pred_str) == len(label_str):
            for index in random.sample(range(len(label_str)), 3):
                print(f'reference: "{label_str[index]}"')
                print(f'predicted: "{pred_str[index]}"')

        else:
            for index in random.sample(range(len(label_str)), 3):
                print(f'reference: "{label_str[index]}"')
                print(f'predicted: "{pred_str}"')

    # Compute WER
    wer = wer_metric.compute(predictions=pred_str, references=label_str)

    # Compute CER
    cer = cer_metric.compute(predictions=pred_str, references=label_str)

    return {"wer": wer, "cer": cer}
```

```
#mmmmiiiiinnnnneeeee
```

```
# from transformers import Wav2Vec2ForCTC
```

```
# model = Wav2Vec2ForCTC.from_pretrained(
#     # "facebook/wav2vec2-xls-r-300m",
#     'facebook/wav2vec2-large-xlsr-53',
#     attention_dropout=0.05,
#     hidden_dropout=0.1,
#     feat_proj_dropout=0.1,
#     mask_time_prob=0.05,
#     layerdrop=0.01377,
#     gradient_checkpointing=True,
#     ctc_loss_reduction="mean",
#     ctc_zero_infinity=True,
#     pad_token_id=processor.tokenizer.pad_token_id,
#     vocab_size=len(processor.tokenizer),
```

```
# )
```

➦ Some weights of Wav2Vec2ForCTC were not initialized from the model checkpoint at facebook/wav2vec2-large-xlsr-53 and are
 You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
#pppeeeerrrrrseeaaain-----iran
```

```
from transformers import Wav2Vec2ForCTC
```

```
model = Wav2Vec2ForCTC.from_pretrained(
    'facebook/wav2vec2-large-xlsr-53',
    attention_dropout=0.05,
```

```

activation_dropout=0.1,
hidden_dropout=0.1,
feat_proj_dropout=0.01249,
final_dropout=0.0,
mask_time_prob=0.05,
mask_time_length=10,
mask_feature_prob=0,
mask_feature_length=10,
layerdrop=0.01377,
gradient_checkpointing=True,
ctc_loss_reduction="mean",
ctc_zero_infinity=True,
bos_token_id=processor.tokenizer.bos_token_id,
eos_token_id=processor.tokenizer.eos_token_id,
pad_token_id=processor.tokenizer.pad_token_id,
vocab_size=len(processor.tokenizer.get_vocab())
)

```

⚠ Some weights of Wav2Vec2ForCTC were not initialized from the model checkpoint at facebook/wav2vec2-large-xlsr-53 and are
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
model.freeze_feature_encoder()
```

```
# import huggingface_hub
```

```
# huggingface_hub.login()
```

```
#repo_name = "wav2vec2-kashmiri-jhon-data-one"
```

```
save_dir = "/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/2training_
```

```
#MMMMMMIIINNNEEEEE
```

```
from transformers import TrainingArguments
```

```

training_args = TrainingArguments(
    output_dir=save_dir,
    group_by_length=True,
    per_device_train_batch_size=8,
    per_device_eval_batch_size=8,
    gradient_accumulation_steps=2,
    evaluation_strategy="steps",
    num_train_epochs=30,
    fp16=True,
    save_steps=500,
    eval_steps=500,
    logging_steps=10,
    learning_rate=4e-4,
    warmup_steps=250,
    save_total_limit=2,
    dataloader_num_workers=24
)

```

⚠ /home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/training_args.py:1594: FutureWarning: `eval`
warnings.warn(

```

# #PERSIAN

# from transformers import TrainingArguments

# training_args = TrainingArguments(
#     output_dir=repo_name,
#     group_by_length=True,
#     per_device_train_batch_size=2,
#     gradient_accumulation_steps=2,
#     eval_strategy="steps",
#     num_train_epochs=20,
#     gradient_checkpointing=True,
#     fp16=True,
#     save_steps=20,
#     eval_steps=20,
#     logging_steps=40,
#     learning_rate=3e-4,
#     warmup_steps=50,
#     save_total_limit=2,
#     push_to_hub=True,
# )

# import numpy as np
# from transformers import Trainer
# trainer = Trainer(
#     model=model,
#     data_collator=data_collator,
#     args=training_args,
#     compute_metrics=compute_metrics,
#     train_dataset=train_dataset,
#     eval_dataset=test_dataset,
#     tokenizer=processor.feature_extractor,
# )

import numpy as np
from transformers import Trainer

# Assuming processor is an instance of Wav2Vec2Processor (or similar for your model)
trainer = Trainer(
    model=model,
    data_collator=data_collator,
    args=training_args,
    compute_metrics=compute_metrics,
    train_dataset=train_dataset,
    eval_dataset=test_dataset,
    processing_class=processor, # Use the processor directly for feature extraction
)

print("step1")
train_result = trainer.train()
print("step2")

metrics = train_result.metrics
print("step3")
max_train_samples = len(train_dataset)
metrics["train_samples"] = min(max_train_samples, len(train_dataset))
print("step4")
trainer.save_model()
print("model created!")
trainer.log_metrics("train", metrics)
trainer.save_metrics("train", metrics)
trainer.save_state()

```

Step	Training Loss	Validation Loss	Wer	Cer
500	2.490800	2.142393	1.000000	0.708172
1000	0.573100	0.633072	0.576649	0.167976

```
> python train.py --num_gpus=8 --local_rank=-1 --master_port=7090 &&
```

```
Trainer.tokenizer is now deprecated. You should use Trainer.processing_class instead.
```

```
2500 /home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/torch/utils/data/dataloader.py:624: UserWarning: This Da
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
# trainer.push_to_hub()
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```

```
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
```

```
warnings.warn(
```



```

import torch
import torchaudio
import librosa
import numpy
from transformers import Wav2Vec2ForCTC, Wav2Vec2Processor
from transformers import Wav2Vec2Processor

/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing wav2vec2.py:174:
# model_name_or_path = "/home/muzaaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment
# device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
# print(model_name_or_path, device)

# processor = Wav2Vec2Processor.from_pretrained(model_name_or_path)
# model = Wav2Vec2ForCTC.from_pretrained(model_name_or_path).to(device)

# def speech_file_to_array_fn(batch):
#     speech_array, sampling_rate = torchaudio.load(batch["file_path"])
#     speech_array = speech_array.squeeze().numpy()
#     #speech_array = librosa.resample(np.asarray(speech_array), sampling_rate, processor.feature_extractor.sampling_rate)
#     speech_array = librosa.resample(y=np.asarray(speech_array), orig_sr=sampling_rate, target_sr=processor.feature_extractor.sampling_rate)

#     batch["speech"] = speech_array
#     return batch

# def predict(batch):
#     features = processor(
#         batch["speech"],
#         sampling_rate=processor.feature_extractor.sampling_rate,
#         return_tensors="pt",
#         padding=True
#     )

#     input_values = features.input_values.to(device)
#     #attention_mask = features.attention_mask.to(device)
#     attention_mask = features.attention_mask.to(device) if "attention_mask" in features else None

#     with torch.no_grad():
#         logits = model(input_values, attention_mask=attention_mask).logits

#     pred_ids = torch.argmax(logits, dim=-1)

#     batch["predicted_N_LM"] = processor.batch_decode(pred_ids)
#     return batch

# import torchaudio
# import librosa
# from datasets import load_dataset
# import numpy as np

# dataset = load_dataset("csv", data_files={"/home/muzaaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment
# dataset = dataset.map(speech_file_to_array_fn)

/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing wav2vec2.py:174:
def speech_file_to_array_fn(batch):
    speech_array, sampling_rate = torchaudio.load(batch["file_path"])
    speech_array = speech_array.squeeze().numpy() # Convert to numpy array

    # Ensure the audio is always a 1D NumPy array (sometimes it's multi-channel)
    if len(speech_array.shape) > 1:
        speech_array = np.mean(speech_array, axis=0) # Convert stereo to mono

    # Resample to match the processor's expected sample rate
    speech_array = librosa.resample(
        y=np.asarray(speech_array),
        orig_sr=sampling_rate,
        target_sr=processor.feature_extractor.sampling_rate
    )

    batch["speech"] = speech_array.tolist() # Convert to Python list (ensures consistency)
    return batch

return batch

```

```

model_name_or_path = "/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/"
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(model_name_or_path, device)

processor = Wav2Vec2Processor.from_pretrained(model_name_or_path)
model = Wav2Vec2ForCTC.from_pretrained(model_name_or_path).to(device)

def predict(batch):
    features = processor(
        batch["speech"],
        sampling_rate=processor.feature_extractor.sampling_rate,
        return_tensors="pt",
        padding=True
    )

    input_values = features.input_values.to(device)
    #attention_mask = features.attention_mask.to(device)
    attention_mask = features.attention_mask.to(device) if "attention_mask" in features else None

    with torch.no_grad():
        logits = model(input_values, attention_mask=attention_mask).logits

    pred_ids = torch.argmax(logits, dim=-1)

    batch["predicted_N_LM"] = processor.batch_decode(pred_ids)
    return batch

import torchaudio
import librosa
from datasets import load_dataset
import numpy as np

dataset = load_dataset("csv", data_files={"/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/
dataset = dataset.map(speech_file_to_array_fn)
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
/home/muzaffar/Desktop/Research/papers/5-paper Wav2Vec/5. Wave2vec Whisper Paper/KASHMIRI/experiment5/2training_all_data
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
warnings.warn(
result = dataset.map(predict, batched=True, batch_size=4)
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
warnings.warn(
from evaluate import load # Use `evaluate` instead of `datasets`

# Load WER and CER metrics
wer = load("wer")
cer = load("cer")

# Compute WER and CER using the correct split
print("WER: {:.2f}".format(100 * wer.compute(predictions=result["train"]["predicted_N_LM"],
                                             references=result["train"]["Transcription"])))
print("CER: {:.2f}".format(100 * cer.compute(predictions=result["train"]["predicted_N_LM"],
                                             references=result["train"]["Transcription"])))

/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
UserWarning: You are not caching version of the module from /home/muzaffar/.cache/huggingface/modules/evaluate_modules/metrics/ev
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
UserWarning: You are not caching version of the module from /home/muzaffar/.cache/huggingface/modules/evaluate_modules/metrics/ev
WER: 13.30
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
warnings.warn(
for i in range(len(result["train"])): # Specify the "train" split
    reference = result["train"]["Transcription"][i] # Use "Text" as reference
    predicted_N_LM = result["train"]["predicted_N_LM"][i]

    if reference.strip() == predicted_N_LM.strip():
        continue

    print("Reference:", reference)
    print("Predicted:", predicted_N_LM)
    print('----')

    predicted: "رنلچوتکرب"
/home/muzaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/torch/utils/data/dataloader.py:624: UserWarning: This Da

```

```

/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
ReferenceWarning: واتناوان ميو ڪُلي زوان تم گوشه نشيني اختيار (
Phode/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
--warnings.warn(
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
Predictions.warn گته ڪري نه پونريو تم دودے سانم کانپ
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
ReferenceWarning: اته وعڊ ڪران و آلي چه سوال ڪران ز تيئلم ڪياز ڪريے پانم
Phode/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
--warnings.warn(
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
Predictions.warn چهس بونی سو فریس سہ ڪته تم ڪجه پنز ب
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
ReferenceWarning: شايم تم آنر ڏنم لولم ڪته بے اوتار بهت پئم شاعر
Phode/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
--warnings.warn(
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
Predictions.warn راڙن باوڪه ناه دوه ميانی تعريقی پيئھ ڪياه
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
ReferenceWarning: تيمن ستي منسوب واقعات چه ڪانس لڪم شاعرن موزون ڪر
Phode/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
--warnings.warn(
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
Predictions.warn چه ڪابلم قبول تم امم خاندان ڪم شجر نسٽڪ
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
ReferenceWarning: ايڇ شوب زون اس تم فقط يوسف شيني محل محل خانج ڪونگ
Phode/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
--warnings.warn(
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
Predictions.warn بسان رديف تم ڪافياھ چهس شعر پت شعر
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
ReferenceWarning: ڪينڙن بند خيال چه ز زول چه پوان زانم شاه صابن
Phode/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
--warnings.warn(
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
Predictions.warn تم سنڌڪ پے تم ڪانہم پتيھ بے
/home/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:
ReferenceWarning: پانم وني سلامه بنديخاه ڪرته ت
Phode/muzaaffar/anaconda3/envs/tf14/lib/python3.11/site-packages/transformers/models/wav2vec2/processing_wav2vec2.py:174:

```